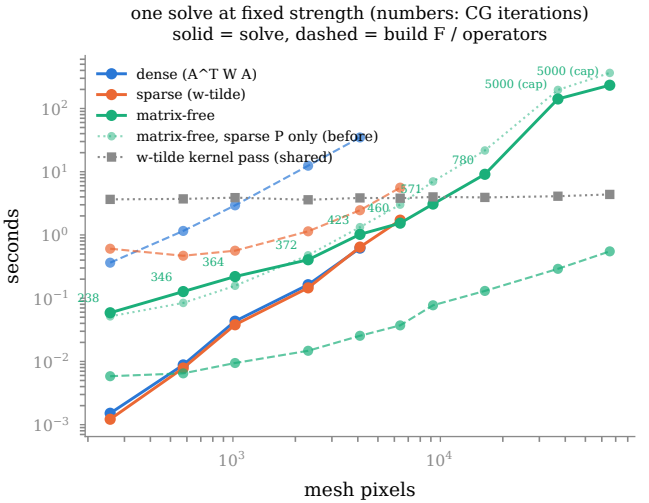
one solve at fixed strength (numbers: CG iterations)
solid = solve, dashed = build F / operators
dense (A^T W A)
sparse (w-tilde)
matrix-free
matrix-free, sparse P only (before)
w-tilde kernel pass (shared)
seconds
mesh pixels
5000 (cap)
5000 (cap)
780
371
460
423
372
364
346
238

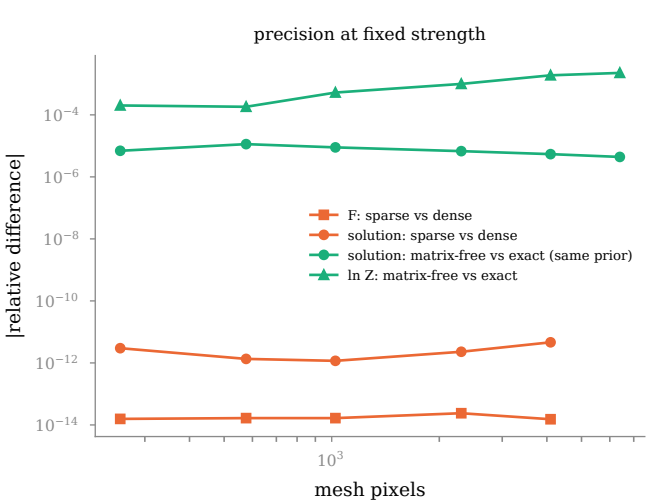
precision at fixed strength
|relative difference|
F: sparse vs dense
solution: sparse vs dense
solution: matrix-free vs exact (same prior)
ln Z: matrix-free vs exact
mesh pixels

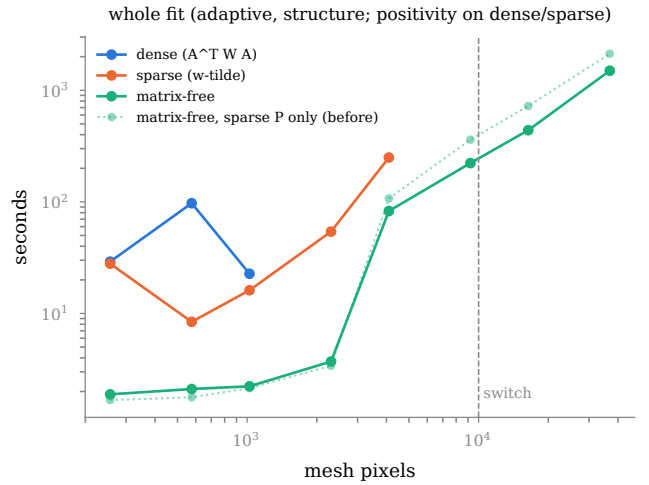
whole fit (adaptive, structure; positivity on dense/sparse)
dense (A^T W A)
sparse (w-tilde)
matrix-free
matrix-free, sparse P only (before)
seconds
switch
mesh pixels

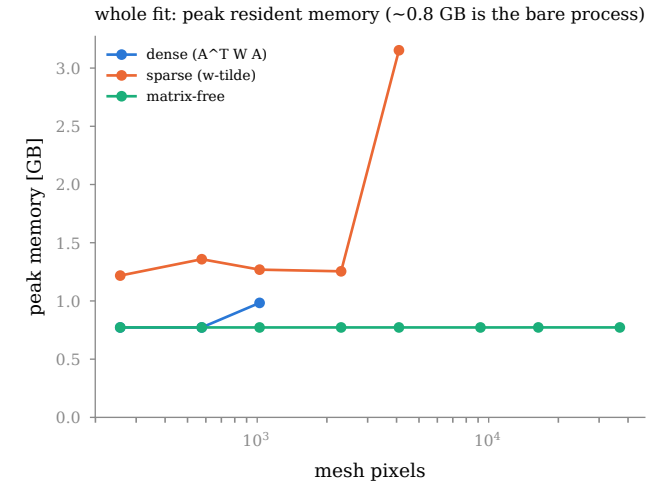
whole fit: peak resident memory (~0.8 GB is the bare process)
dense (A^T W A)
sparse (w-tilde)
matrix-free
peak memory [GB]
mesh pixels